

Based on
Mastering Networks - An Internet Lab Manual
by Jörg Liebeherr and Magda Al Zarki

Adapted for
'Labo Computernetwerken'
by Johan Bergs, Nicolas Letor, Michael Voorhaen and Kurt Smolderen

Completed by
Josse Coen Armin Halilovic Jonas Vanden Branden Group 2

May 19, 2016

Lab 6

LAN switching

What you will learn in this lab:

- How to configure a Cisco Router and a Linux PC as a LAN switch.
- How LAN switches update their forwarding tables.
- How LAN switches run a spanning tree protocol for loop free routing.

6.1 Prelab 6

Bridges and the Spanning Tree Protocol

Use the following resources to prepare yourself for this lab session:

1. Bridging: Read about LAN switching and bridging at http://docwiki.cisco.com/wiki/Internetworking_Technology_Handbook#Bridging_and_Switching.
2. Transparent Bridges and Spanning Tree Protocol: Read about transparent bridges and the spanning tree protocol at http://docwiki.cisco.com/wiki/Transparent_Bridging.
3. Bridge Protocol Data Unit (BPDU): Familiarize yourself with the format of bridge protocol data units (BPDUs) by reading the information at <http://ericleahy.com/index.php/implementing-spanning-tree-protocol-stp/>.
4. Configuring a PC as a Bridge: Explore the website <http://www.linuxfoundation.org/collaborate/workgroups/networking/bridge>, which describes the bridge-utils software package for configuring a Linux PC as a bridge.

Prelab Questions

Question 1)

Describe the difference between a LAN switch/bridge and a router?

A switch/bridge connects different hosts to form a network. It works on the link layer, so it only looks up to the Ethernet header of a packet. Switches/bridges never analyze IP addresses, for instance.

A router works on the transport layer: It does IP routing, and connects different IP networks to each other.

Question 2)

What is the difference between an Ethernet switch and an Ethernet hub? Which is more suitable for a network with a high traffic load, a switch or a hub? Explain.

A hub sends all incoming packets to all its outgoing interfaces and the host/router who can do something with the packets will handle them. A switch maps hardware addresses with one of its interfaces and checks to where a packet should be sent. A switch is clearly more suitable for a high traffic network, as the hub will pollute all of its interfaces with duplicates of every packet, causing collisions.

Question 3)

What motivates the use of the term "transparent" in transparent bridges?

They are called transparent because their operation is transparent to the hosts they are connected to. It learns which hardware addresses are to be mapped to which interfaces just through the passing of packets on their way, never having to explicitly communicate to the hosts about wanting to learn where they are. If the bridge receives a packet with a destination hardware address it doesn't know yet, it just forwards it to all interfaces, like a hub does. Once the unknown host sends a packet of its own, the bridge saves its interface, and packets destined for that host won't have to be sent out to all interfaces anymore.

Question 4)

Which role does the spanning tree protocol play when interconnecting LAN switches/bridges?

The spanning tree protocol (STP) is used to control redundancy of paths in the network topology. If a switch using STP detects that packets are being sent to the same network segment multiple times, one of the paths can be broken by turning off the port. In fact, it will prevent any loops from existing within the network by determining a spanning tree (a graph without any cycles).

Question 5.a)

In the context of the IEEE 802.1d specification of the spanning tree protocol, define root bridge.

A root bridge is a bridge selected to be the root of the spanning tree that is created, from which the tree will branch out to other bridges.

Question 5.b)

In the context of the IEEE 802.1d specification of the spanning tree protocol, define root port.

The root port of a bridge is a bridge's port with least path cost to the root bridge. A root bridge does not have a root port, while all other bridges have exactly one root port.

Question 5.c)

In the context of the IEEE 802.1d specification of the spanning tree protocol, define designated bridge.

The designated bridge of a LAN segment is the bridge that provides the minimum root path cost. While a bridge only has one root port, it can have multiple designated ports. All the

ports on a root bridge are designated ports. A port can't be a root port and a designated port at the same time.

Question 5.d)

In the context of the IEEE 802.1d specification of the spanning tree protocol, define designated port.

A port on a bridge is called designated if the bridge is designated on the LAN segment the port connects to.

Question 5.e)

In the context of the IEEE 802.1d specification of the spanning tree protocol, define blocked port.

As explained in question 4, paths that form cycles will be stopped from existing. Once a bridge determines that a cycle exists, it will turn off one of the ports that form the cycle, breaking the cycle. Such a port is called a blocked port. A port in blocking state will never send any frames through the port, and discard any packets it receives from the port.

Question 6)

In the spanning tree protocol, how does a LAN switch/bridge decide which ports are in a blocking state?

For every path, the path cost to the root bridge is determined. Any path with a cost greater than the path with the minimum path cost will be blocked. Obviously, the root port of a bridge will never be blocked. Also, designated ports won't be blocked.

6.2 Lab 6

A bridge or LAN switch is a device that interconnects two or more Local Area Networks (LANs) and forwards packets between these networks. Different from IP routers, bridges and LAN switches operate at the data link layer. For example, bridges and LAN switches forward packets based on MAC addresses, whereas IP routers forward packets based on IP addresses.

LAN switches are widely deployed in enterprise networks, including university campus networks. Many enterprise networks primarily use LAN switches, and use IP routers only to connect the enterprise network to the public Internet.

The term 'bridge' was coined in the early 1980s. Today, when referring to data-link layer interconnection devices, the terms 'LAN switch' or 'Ethernet switch' (in the context of Ethernet) are much more common. Since many of the concepts, configuration commands, and protocols for LAN switches in Lab 6 use the old term 'bridge', we will, with few exceptions, refer to LAN switches as bridges.

This lab covers the main concepts of LAN switching in Ethernet networks: how packets are forwarded between LANs and how the routes of packets are determined. In the first and second parts of Lab 6, you learn how to configure a Linux PC and a Cisco router as a bridge. The third part illustrates the difference between an Ethernet hub and an Ethernet switch. Part 4 explores how forwarding tables of bridges are set up. You learn about the concepts of learning bridges and transparent bridges.

The configuration of the equipment in Lab 6 is changed several times during the course of the lab. With exception of the last part, the IP address configuration of the Linux PCs is as shown in Table ???. Note that all IP addresses have the same netmask.

Linux PC	eth0	eth1
PC1	10.0.1.11/24	10.0.1.12/24
PC2	10.0.2.21/24	10.0.1.22/24
PC3	10.0.3.31/24	10.0.1.32/24
PC4	10.0.4.41/24	10.0.1.42/24

Table 6.1: IP addresses

Part 1. Configuring a Linux PC as a bridge

The exercises in this lab show how to configure a Linux PC as a bridge. Ethernet bridging functionality is integrated in all recent versions of Linux. The configuration of bridging functions in Linux is done with configuration commands and tools. In this lab, we use the command-line bridge configuration tool `brctl`.

The network configuration for this part is shown in Figure ???. Here, PC1 and PC3 act as hosts and PC2 is set up as a bridge.




Figure 6.1: Network Topology for Part 1.

Exercise 1-a. IP configuration of Linux PCs

1. Set up the network configuration as shown in Figure ??.
2. Configure the interfaces of PC1, PC2, and PC3, with the IP addresses given in Table ??. Disable the interfaces that are not used in the configuration, that is, disable interface *eth1* on both PC1 and PC3.
3. Since, throughout Lab 6, you frequently work with MAC addresses, you should record the MAC addresses of the Linux PCs. Log in to each of the PCs and obtain the MAC addresses of both Ethernet interfaces with the command `ifconfig -a`. Enter the MAC addresses in Table ??. (In Part 2, you will repeat the same exercise on the Cisco routers.)

Linux PC	MAC address of eth0	MAC address of eth1
PC1	68:05:ca:36:33:a0	68:05:ca:39:cc:79
PC2	68:05:ca:36:31:f0	68:05:ca:39:e1:36
PC3	68:05:ca:36:39:c7	68:05:ca:36:51:3f
PC4	68:05:ca:39:e1:2f	68:05:ca:39:e1:32

Table 6.2: MAC addresses of the linux PCs

Exercise 1-b. Configure a Linux PC as a bridge

In this exercise, you configure PC2 as a bridge that forwards packets between the two Ethernet segments shown in Figure ??. The bridge configuration on the Linux PCs is done with the tool `brctl`.

1. Creating a bridge with `brctl`: It is possible to configure multiple independently operating bridges on the same PC. Each bridge is assigned a name and is associated with a set of interfaces. Here, you configure one bridge on PC2 and assign the bridge the name `Bridge1`.

```
| brctl addbr Bridge1
```

Check if the bridge has been created by typing:

```
| brctl show
```

2. Configuring a bridge with `brctl`: After the bridge is created, the bridge is configured in the following steps:

- Assign interfaces to the bridge. e.g. for PC2, add the interfaces `eth0` and `eth1`.

```
| brctl addif Bridge1 eth0
| brctl addif Bridge1 eth1
```

- The next part of the configuration sets the parameters of the spanning tree protocol (STP). In Part 1, the spanning tree protocol is not used. Therefore, you need to disable the spanning tree protocol by toggling the button next to the STP label, so that it shows the label Disabled.

```
| brctl stp Bridge1 off
```

- In the last part of the bridge configuration, you activate the bridge `Bridge1` from a terminal window. On a Linux PC, each created bridge is represented as a network interface. Therefore, if you type the command `ifconfig -a` on PC2, the command shows an interface `Bridge1`, in addition to the other interfaces `eth0`, `eth1`, and `lo`. The bridge is activated by enabling the interface associated with the bridge. This is done with the following command:

```
| PC1% ifconfig Bridge1 up
```



Activating the bridge disables the IP configuration of the interfaces assigned to a bridge. Hence, it is no longer possible to issue ping commands to these interfaces.



The following settings can also be configured on a bridge interface. You don't need to change them for this exercise however.

```
setageing <bridge> <time>
```

Set ageing time to <time> for bridge interface <bridge>

```
setbridgeprio <bridge> <prio>
```

Set bridge priority to <prio> for bridge interface <bridge>

```
setfd <bridge> <time>
```

Set bridge forward delay to <time> for bridge interface <bridge>

```
sethello <bridge> <time>
```

Set hello time to <time> for bridge interface <bridge>

```
setmaxage <bridge> <time>
```

Set the maximum message age to <time> for bridge interface <bridge>

```
stp <bridge> on|off
```

Turn the Spanning Tree Protocol on/off for bridge interface <bridge>

Exercise 1-c. Observing a bridge in operation

You will use Wireshark in this exercise. Do not forget to append the binary dump (pcap format) to your lab report

When the bridge configuration of PC2 is complete, PC2 forwards packets between PC1 and PC3. This exercise asks you to observe the forwarding.

1. Start Wireshark on PC1 and PC3, and capture traffic on interface *eth0* on both systems.
2. When bridging is activated on PC2, the configured IP addresses on PC2 should be disabled. To verify this, issue a ping command to interfaces *eth0* and *eth1* of PC2 from PC1 and PC3.

```
PC1% ping 10.0.1.21
PC3% ping 10.0.1.22
```

If PC2 is configured as a bridge, these ping commands should fail.

3. Clear the ARP caches on PC1 and PC3. Note that in Linux, each ARP entry has to be deleted separately with the command `arp -d <MACaddress>`.
4. Issue a ping command from PC1 to PC3 and save the output.

```
PC1% ping -c 1 10.0.1.31
```

Observe that PC2 actually forwards the packets between PC1 and PC3.

Question 1.C.1)

Do the source and destination MAC/IP addresses change when a packet traverses a bridge? Provide an explanation and include an example from the captured data. Suppose that PC2 was configured as an IP router, which differences would you observe in the Ethernet and IP headers?

No, the source and destination MAC/IP addresses do not change when a packet traverses a bridge. A bridge does not have an address itself; it only forwards packets that are on their way to objects that do have an address, such as routers and hosts.

An example:

PC1

No.	Time	Source	Destination	Protocol	Length
1	7.79.726862	10.0.1.11	10.0.1.31	ICMP	98
3	Echo (ping) request id=0x0c87, seq=1/256, ttl=64 (reply in 8)				
5	Frame 7: 98 bytes on wire (784 bits), 98 bytes captured (784 bits)				
7	Ethernet II, Src: IntelCor_36:33:a0 (68:05:ca:36:33:a0), Dst: IntelCor_36:39:c7 (68:05:ca:36:39:c7)				
9	Internet Protocol Version 4, Src: 10.0.1.11 (10.0.1.11), Dst: 10.0.1.31 (10.0.1.31)				
11	Version: 4				
13	Header length: 20 bytes				
15	Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))				
	Total Length: 84				
	Identification: 0xf239 (62009)				
	Flags: 0x02 (Don't Fragment)				
	Fragment offset: 0				
	Time to live: 64				
	Protocol: ICMP (1)				

```

17   Header checksum: 0x3246 [validation disabled]
    Source: 10.0.1.11 (10.0.1.11)
    Destination: 10.0.1.31 (10.0.1.31)
19   [Source GeoIP: Unknown]
    [Destination GeoIP: Unknown]
21 Internet Control Message Protocol

```

PC3

```

1  No.      Time      Source      Destination      Protocol Length
   Info
   11 79.727330 10.0.1.11    10.0.1.31      ICMP      98
   Echo (ping) request id=0x0c87, seq=1/256, ttl=64 (reply in 12)
3
Frame 11: 98 bytes on wire (784 bits), 98 bytes captured (784 bits)
5 Ethernet II, Src: IntelCor_36:33:a0 (68:05:ca:36:33:a0), Dst: IntelCor_36:39:c7
  (68:05:ca:36:39:c7)
Internet Protocol Version 4, Src: 10.0.1.11 (10.0.1.11), Dst: 10.0.1.31 (10.0.1.31)
7  Version: 4
   Header length: 20 bytes
9   Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (
     Not ECN-Capable Transport))
   Total Length: 84
11  Identification: 0xf239 (62009)
   Flags: 0x02 (Don't Fragment)
13  Fragment offset: 0
   Time to live: 64
15  Protocol: ICMP (1)
   Header checksum: 0x3246 [validation disabled]
17  Source: 10.0.1.11 (10.0.1.11)
   Destination: 10.0.1.31 (10.0.1.31)
19  [Source GeoIP: Unknown]
   [Destination GeoIP: Unknown]
21 Internet Control Message Protocol

```

We see here that nothing was changed.

If PC2 was configured as an IP router, then the IP addresses of the packets would remain the same (given that no NAT was configured on PC2), but the destination field of the packet traveling towards PC2 would be PC2's hardware address instead of PC3's hardware address, and the forwarded packet would have PC2's hardware address in the source field instead of PC1's hardware address.

- Run traceroute from PC1 to PC3 and save the output

```
PC1% traceroute 10.0.1.31
```

Here, you should observe that PC2 does not appear in the output of traceroute.

Question 1.C.2)

Include the output from the traceroute command. Why is PC2 not visible from PC1?

```

1 student@lab2pc1:~$ traceroute 10.0.1.31
  traceroute to 10.0.1.31 (10.0.1.31), 30 hops max, 60 byte packets
3  1  10.0.1.31 (10.0.1.31)  2.567 ms  2.872 ms  2.869 ms

```

The way traceroute works is by causing a time-to-live timeout at every intermediary IP hop. The router where such a timeout happens sends an ICMP error packet back to the originator of the packet that timed out. This is how the source knows about the hop. A bridge does not decrease the TTL field of a packet, nor does it have an address from which it can send packets (as discussed earlier).

Question 1.C.3)

If PC2 was configured as an IP router, how would the output differ?

If PC2 was configured as an IP router, it would be an IP hop with its own hardware addresses and IP addresses, and it would decrease the TTL of packets. As such, the TTL of the packet would be able to exceed at PC2, and PC2 would be able to send an ICMP TTL exceeded error message. Upon receiving this error message, PC1 would know of PC2's existence, and it would be added as hop on traceroute.

6. Change the IP address of PC3 to 10.0.2.12/24. Note that PC1 and PC3 now have different IP network prefixes. Repeat Step 4.

Question 1.C.4)

Does the ping command from PC1 to PC3 still work? Explain the outcome.

The ping command no longer works. This happens because the hosts are not on the same subnet anymore. Every host/router connected to a bridge on a particular interface is supposed to be on the same IP network on that interface. Because PC1 doesn't know about the network PC3 is on, and doesn't have a default gateway either, it doesn't know where to send the ping message.

Exercise 1-d. Manipulating a PC bridge

This exercise familiarizes you with a few tasks related to running `brctl` on a Linux PC. You learn how to display the MAC forwarding table, how to delete the contents of the MAC forwarding table, and, finally, how to turn off the bridging functions. All of the tasks are performed on PC2.

1. First, reset the IP address of the `eth0` interface of PC3 to 10.0.1.31/24.
2. Displaying the MAC forwarding table: The MAC forwarding table of a bridge plays the same role as the routing table of an IP router. To view the contents of the MAC forwarding table of Bridge1 on PC2, perform the following steps:

```
| brctl showmacs Bridge1
```

3. Clearing the MAC forwarding table of a bridge: The `brctl` tool does not have a convenient way to delete the contents of the MAC forwarding table. Instead you must exploit that a bridge automatically deletes an entry in the forwarding table that has not been looked up for a certain time, which is determined by the ageing parameter of `brctl`. To delete the entries in the forwarding table, you must set the ageing parameter to zero seconds. Once the entries are deleted, set the Ageing entry back to the original value (The default value is 300 seconds).

```
| brctl setageing Bridge1 0
| brctl setageing Bridge1 300
```

4. Disabling a bridge: Disabling a bridge on a Linux PC is done in two steps: (1) deactivate the interface associated with `Bridge1` and (2) delete the bridge.

```
| ifconfig Bridge1 down  
| brctl delbr Bridge1
```

You can verify that the bridge is disabled as follows: Verify that PC2 is operating as a normal host. To do this, issue a ping command to interfaces *eth0* and *eth1* of PC2 from PC1 and PC3.

```
| PC1% ping -c 1 10.0.1.21  
| PC3% ping -c 1 10.0.1.22
```

Verify that PC2 does not forward packets by issuing the following ping command:

```
| PC1% ping -c 1 10.0.1.31
```

The ping command should not be successful.

Part 2. Configuring a Cisco Router as a bridge

Next you learn how to configure a Cisco Router as a bridge. The topology for this part is shown in Figure ???. Router1 is configured as a bridge that connects the two Ethernet segments.




Figure 6.2: Network Topology for Parts 2 and 3-b.

Exercise 2-a. Setup of network configuration

After the network is configured, as in Exercise 1-a above for the Linux PCs, you are asked to record the MAC addresses of the Cisco routers.

1. Connect the PCs and Router1 with Ethernet hubs as shown in Figure ??.
2. Configure the *eth0* interfaces of the Linux PCs with the IP addresses given in Table ??. (The IP addresses of PC1, PC2, and PC3 are the same as in Part 1). Disable the *eth1* interfaces of all Linux PCs.
3. Establish a *minimcom* session to each Cisco router: from PC1 to Router1, from PC2 to Router2, and so on.
4. On each router, type *enable* and then use the command *show interfaces* to display the MAC addresses of the Ethernet interfaces. Record the MAC addresses in Table ??

Linux PC	MAC address of FastEthernet0/0	MAC address of interface vlan1
Router1	000d.6517.0129	0012.00d4.bf58
Router2	000d.bcef.b1a3	0012.0120.245c
Router3	000d.bcef.eb24	0012.00d4.bdb0
Router4	000d.bcef.b948	0012.00d4.bf6c

Table 6.3: MAC addresses of the linux PCs

Exercise 2-b. Configuring a Cisco Router as a bridge

Next you configure Router1 as a bridge. A Cisco router is configured as a bridge by disabling IP forwarding functions (with the command `no ip routing`) and by enabling bridging functions.

Similar as on the Linux PCs, a Cisco router can be configured to perform the functions of multiple independently operating bridges. This is done by defining a bridge group, which is identified by a number, and associating two or more network interfaces with each bridge group. Packets are forwarded only between interfaces that are assigned to the same bridge group. Since the exercises in Lab 6 only use one bridge group, we always use 1 to identify the group.

In order to create a bridge on the Cisco router, go into IOS's Global Configuration Mode and execute the following commands:

```
Router1(config)# bridge 1 protocol ieee
Router1(config)# bridge 1 priority 128
```

The first command defines a bridge group identified by number 1 and assigns the spanning tree protocol as defined in IEEE 802.1d to bridge group 1. After the command is issued, the Cisco router forwards packets between all interfaces that are assigned to bridge group 1. A bridge group can be any number between 1 and 63. After defining a bridge group, one can assign network interfaces to the bridge group. It is possible to define multiple bridge groups. In Lab 6, only one bridge group (with identifier 1) is used. The second command assigns the priority 128 to bridge group 1. The priority of a bridge group plays a role in the spanning tree protocol, which is covered in Part 5.



Each interface is individually configured to participate in a bridge group. This is done with the following commands (all need to be run from the IOS Interface Configuration Mode):

```
bridge-group 1
```

Assign the current network interface to bridge group 1.

```
no bridge-group 1
```

Remove the current network interface from bridge group 1.

```
bridge-group 1 spanning-disabled
```

Disable the spanning tree protocol on the current interface for bridge group 1.

```
no bridge-group 1 spanning-disabled
```

Enable the spanning tree protocol on the current interface for bridge group 1.



Once a Cisco router is configured as a bridge, the following commands can be used to display the status of the bridge (all need to be run from the Privileged Exec IOS Mode):

```
show bridge
```

Display the entries of the MAC forwarding table.

```
show spanning-tree
```

Display the spanning tree topology information known to this bridge.

```
show interfaces
```

Display statistics of all interfaces, including the MAC addresses of all interfaces.



The following commands disable bridging functions on a Cisco router. The first must be executed in IOS's Global Configuration Mode, the subsequent commands must be run from the Privileged Exec IOS Mode):

```
no bridge 1
```

Delete the defined bridge group. After the command is issued, the Cisco router stops forwarding packets between interfaces that are assigned to bridge group 1.

```
clear bridge
```

Remove all entries from the MAC forwarding table.

```
clear arp-cache
```

Clear the ARP table.

1. Configure a Cisco Router as a Bridge: Use the above commands to configure Router1 as a bridge. On Router1, type the following commands:

```
Router1> enable Password: <enable secret>
Router1# configure terminal
Router1(config)# no ip routing
Router1(config)# bridge 1 protocol ieee
Router1(config)# bridge 1 priority 128
Router1(config)# interface FastEthernet0/0
Router1(config-if)# bridge-group 1
Router1(config-if)# bridge-group 1 spanning-disabled
Router1(config-if)# no shutdown
Router1(config-if)# interface FastEthernet0/1
Router1(config-if)# no shutdown
Router1(config-if)# interface vlan1
Router1(config-if)# bridge-group 1
Router1(config-if)# bridge-group 1 spanning-disabled
Router1(config-if)# no shutdown
Router1(config-if)# end
Router1# clear bridge
Router1# clear arp-cache
```

The commands disable IP forwarding and set up Router1 as a bridge that runs with priority 128. Both Ethernet interfaces are assigned to the bridge, but the spanning tree protocol is disabled.

2. Once Router1 has been configured as a bridge, configure the Linux PCs as shown in Figure ?? with the IP addresses of Table ??.
3. Delete all entries in the ARP caches of PC1 and PC3.
4. Issue a ping command from PC1 to PC3.

```
PC1% ping -c 1 10.0.1.31
```

5. Run traceroute from PC1 to PC3 and save the output.

```
PC1% traceroute 10.0.1.31
```

Question 2.1)

Include the output of the traceroute command.


```
1 student@lab2pc1:~$ traceroute 10.0.1.31
   traceroute to 10.0.1.31 (10.0.1.31), 30 hops max, 60 byte packets
3  1  10.0.1.31 (10.0.1.31)  1.474 ms  1.668 ms  1.664 ms
```

Question 2.2)

Compare the results to the outcome of the `traceroute` command in Exercise 1-c.

We see that the outcome of the `traceroute` command is almost exactly the same as the output of exercise 1-c. The only difference lies in the timings, they are lower in this exercise.

Question 2.3)

Why is it not possible to issue a `ping` command to Router1?

It is not possible to issue the `ping` command to Router1 because the router is set up to work on the data link layer, so it will not respond to `ping` commands.

The Difference Between an Ethernet Hub and an Ethernet Switch

In this part of the lab, you try to observe the difference between the operation of an Ethernet hub and an Ethernet switch. The main observation to be made is that traffic going over an Ethernet hub may experience collisions, whereas an Ethernet switch does not have collisions.

An Ethernet hub is a relatively simple device that merely repeats a signal received on one network interface (port) to all other ports. When multiple devices connected to the same hub transmit a packet at the same time, the transmissions are corrupted. This is referred to as a collision.

An Ethernet switch, which performs the functions of a bridge for Ethernet segments, is a store-and-forward device. When a packet is received, the Ethernet switch looks up the destination MAC address in its MAC forwarding table, and then forwards the packet to one of its ports. Transmissions on an outgoing link are done one packet at a time, and packets are buffered if multiple packets must be forwarded on the same output port at the same time.

In this context, dual-speed Ethernet hubs, which connect both 10 Mbps (10BaseT) and 100 Mbps (100BaseTX) Ethernet devices, are a special case. A dual-speed hub operates as two Ethernet hubs, one running at 10 Mbps and one running at 100 Mbps, that are connected by a bridge. Thus, there can be collisions between devices that operate at the same speed, but there are no collisions between devices at different speeds.

We point out that it is not always possible to observe collisions on an Ethernet hub. Not only does the rate of collision depend on the traffic load and pattern. In addition, hubs increasingly use internal buffering and avoid collisions in many cases.

Exercise 3-a. Observe collisions on an Ethernet hub

Try to generate collisions by flooding an Ethernet hub with traffic. The network configuration is as shown in Figure ???. You intentionally flood the hub that connects PC1 and PC2 with traffic and hopefully force collisions to occur.




Figure 6.3: Network Topology for Part 3

1. Configure the network as shown in Figure ???. Starting from the network configuration from Part 2 (in Figure ??), disconnect Router1 and connect the two hubs directly with a crossover Ethernet cable (black cable, red RJ45 plugs).

2. Determine the number of collisions on interface *eth0* of PC1 and PC3, that have occurred since the PCs have been rebooted the last time, by typing

```
PC1% ifconfig -a
PC3% ifconfig -a
```

Save the output.

3. Flood the network by generating a large number of ICMP Echo Request and Reply packets between PC1 and PC2, by typing

```
PC1% ping -f 10.0.1.21
```

4. While the above ping command is running, start sending 100 ICMP Echo Request packets from PC3 to PC4.

```
PC3% ping -c 100 10.0.1.41
```

Hopefully, these ping messages cause some collisions at the Ethernet hub.

5. On PC1 and PC3, save the output of `ifconfig -a` again. Observe the number of new collisions that the above experiment has generated.

Question 3.A)

Calculate the number of new collision as seen by PC1 and PC3, in the above exercise. Briefly explain what causes the collisions.

PC1: 7 collisions

PC3: 12 collisions

Collisions happen when hosts try to send data simultaneously over the same shared channel. The hub takes the signal and repeats it over the other links it is connected to. If data is already being sent on one of those links, collisions occur, because the hub will just repeat both signals at the same time.

Exercise 3-b. No collisions when using an Ethernet switch

Repeat the steps of the previous exercise, but place an Ethernet switch (Router1, which has been configured as a bridge), between the two Ethernet hubs.

1. Reconstitute the network configuration shown in Figure ??, by connecting Router1 between the two hubs.
2. Obtain the number of collisions that have been observed on interface *eth0* of PC1 and PC3 by typing

```
PC1% ifconfig -a
PC3% ifconfig -a
```

Save the output.

3. Flood the network with traffic by generating a large number of ICMP Echo Request and Reply packets between PC1 to PC2, by typing

```
PC1% ping -f 10.0.1.21
```

4. Now issue 100 ICMP Echo Request packets from PC3 to PC4:

```
| PC3% ping -c 100 10.0.1.41
```

5. Once again, save the output of the command `ifconfig -a` on PC1 and PC3, and record the number of collisions on interface *eth0*.
6. Access Router1 and run the `show bridge` command to display the bridge forwarding table. Save the data.

Question 3.B)

Use the `ifconfig -a` output to calculate the new collisions seen at the interfaces of PC1 and PC3. Explain the differences between the outcomes in Exercise 3-a and Exercise 3-b.

We see no new collisions this time.

A switch is smarter than a hub, in that it stores packets it receives: it needs to analyze the packet in order to decide where to forward it. Thus it will never dumbly send data received from two different hosts at the same time on the same channel.

Part 3. Learning Bridges

Each bridge has a MAC forwarding table that determines the port where a packet is transmitted from. When a packet arrives, the bridge looks up the destination MAC address of the packet in its MAC forwarding table, and retrieves the outgoing port for this packet. If the destination MAC address is not found in the forwarding table, the bridge floods the packet on all ports, with exception of the port where the packet arrived.

Bridges update their MAC forwarding table using what is called a learning algorithm, which works as follows. A bridge examines the source MAC address of each packet that arrives on a particular port, and memorizes that the source address is reachable via that port. This is done by adding the source MAC address and the port to the MAC forwarding table. The next time the bridge receives a packet which has this MAC address as destination, the bridge finds the outgoing port in its forwarding table. Bridges that run this algorithm are referred to as learning bridges. All currently deployed Ethernet switches execute the learning algorithm.

An entry in the MAC forwarding table is deleted if it is not used (looked up) for a certain amount of time. The maximum time that a MAC address can stay in the forwarding table without a lookup is determined by the *Ageing* value, which is a configuration parameter.

Here you investigate the learning algorithm of bridges. The network configuration is as shown in Figure ??.




Figure 6.4: Network Topology for Part 4

Exercise 4-a. Exploring the learning algorithm of bridges

In this exercise you study how bridges set up their MAC forwarding tables from the network traffic.

1. Set up the network configuration as shown in Figure ??.
2. Establish a `minicom` session to Router1, Router2, and Router3.
 - Configure Router1, Router2, and Router3 as bridges (Disable the spanning tree protocol).
 - On each of the bridges, delete the contents of the MAC forwarding table with the `clear bridge` command.

3. Verify that PC2 is not running as a bridge. If necessary, follow the instructions in Exercise 1-d to disable the bridging functions on PC2. Also, verify that on each PC only interface *eth0* is enabled.
4. Start to capture traffic with *ethereal* on the *eth0* interfaces of PC1, PC2, PC3, and PC4.
5. Clear the ARP cache on PC1, PC2, and PC3.
6. Now, issue a set of *ping* commands. After each command, save the MAC forwarding table on all bridges with the command *show bridge*, and observe how far the ICMP Echo Request and Reply packets travel.

```
PC1% ping -c 1 10.0.1.21
PC2% ping -c 1 10.0.1.11
PC2% ping -c 1 10.0.1.41
PC3% ping -c 1 10.0.1.21
```

7. Stop the traffic capture on the PCs, and save the *ethereal* output.

Question 4.A.1)

Use the captured data to illustrate the algorithm used by bridges to forward packets.

We should note that the “hubs” that connect the PCs to the routers were actually switches instead of hubs when we did the lab, so they were also learning.

When the ping command from PC1 to PC2 is executed, PC1 sends out an ARP request for PC2's IP address. Router1 forwards the ARP request and maps PC1's hardware address to its FastEthernet0/0 interface. The switch (supposed to be hub) that connects Router1, Router2 and PC2 doesn't know PC2 yet, so it forwards the ARP request to both PC2 and Router2. Router2 doesn't know PC2 either, so it forwards the request as well. The same goes for the switch that connects Router2, Router3 and PC3. Indeed: in PC3's trace file we see the ARP request from PC1 asking for PC2's hardware address.

```
1 Frame 13: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)
  Ethernet II, Src: IntelCor_36:33:a0 (68:05:ca:36:33:a0), Dst: Broadcast (ff:ff:ff:
    ff:ff:ff)
3     Destination: Broadcast (ff:ff:ff:ff:ff:ff)
   Source: IntelCor_36:33:a0 (68:05:ca:36:33:a0)
5     Type: ARP (0x0806)
   Padding: 00000000000000000000000000000000
7 Address Resolution Protocol (request)
   Hardware type: Ethernet (1)
9   Protocol type: IP (0x0800)
   Hardware size: 6
11  Protocol size: 4
   Opcode: request (1)
13  Sender MAC address: IntelCor_36:33:a0 (68:05:ca:36:33:a0)
   Sender IP address: 10.0.1.11 (10.0.1.11)
15  Target MAC address: 00:00:00_00:00:00 (00:00:00:00:00:00)
   Target IP address: 10.0.1.21 (10.0.1.21)
```

When PC2 responds to PC1's ARP request, Router1, as well as the switches in between, already know about PC1, so they only forward the response to the right interfaces.

Upon learning PC2's hardware address, PC1 sends its ping message to PC2, and PC2 responds. The ping succeeds.

Five seconds later we see an ARP request from PC2 asking for PC1's hardware address, even though no ping message was ready to be sent at that point. We can't quite explain this request, as PC2 learned PC1's hardware address through PC1's ARP request. Perhaps it's an automatic refresh.

Then, PC2 pings PC1. Again, PC1 and PC2 are already known to all the switches in between. Then, PC2 pings PC4. Router2 and Router3, as well as the switches in between, don't know PC4 yet, so they all forward the request to all interfaces. Thus, even PC1 will receive this request:

```

Frame 16: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)
2 Ethernet II, Src: IntelCor_36:31:f0 (68:05:ca:36:31:f0), Dst: Broadcast (ff:ff:ff:
   ff:ff:ff)
   Destination: Broadcast (ff:ff:ff:ff:ff:ff)
4   Source: IntelCor_36:31:f0 (68:05:ca:36:31:f0)
   Type: ARP (0x0806)
6   Padding: 0000000000000000000000000000000000000000000000000000000000000000
   Address Resolution Protocol (request)
8   Hardware type: Ethernet (1)
   Protocol type: IP (0x0800)
10  Hardware size: 6
   Protocol size: 4
12  Opcode: request (1)
   Sender MAC address: IntelCor_36:31:f0 (68:05:ca:36:31:f0)
14  Sender IP address: 10.0.1.21 (10.0.1.21)
   Target MAC address: 00:00:00_00:00:00 (00:00:00:00:00:00)
16  Target IP address: 10.0.1.41 (10.0.1.41)

```

Lastly, PC3 pings PC2. PC2 is already known to all bridges involved, so the ping request is only forwarded by Router2. When PC3 sends its ARP request asking for PC2's hardware address (again, we don't know why PC3 doesn't just derive this address from PC2's ping packet), all of the bridges get to know PC3: we can find PC3's ARP request in the traces of all PCs.

Question 4.A.2)

For each of the transmitted packets, explain if the learning algorithm results in changes to the MAC forwarding table. Describe the changes.

All of the changes were already discussed in the previous question.

Exercise 4-b. Learning about new locations of hosts

Learning bridges adapt their MAC forwarding tables automatically when the location of a host changes. Due to the learning algorithm, the time it takes to adapt to a change depends on the network traffic and on the value of the Ageing parameter. This is illustrated in the following exercise.

1. Continue with the configuration of the previous exercise. First, create or refresh entries in the MAC forwarding table at the bridges by issuing the following commands from PC1:

```

PC1% ping -c 3 10.0.1.31
PC1% ping -c 3 10.0.1.41

```

2. Now, connect PC2 to the same hub that PC4 is connected to.
3. Issue a ping command from PC1 that continuously sends ICMP Echo Request packets to PC2

```

PC1% ping 10.0.1.21

```

Since Router2 does not know that PC2 has moved, it does not forward the ICMP Echo Request packet, and packet does not reach PC2. As a result, the ARP requests and the

ping are unsuccessful. Eventually, since the MAC forwarding entry for PC2 is not refreshed at Router2 and Router3, the entry is deleted. When the entry is removed, the next ICMP Echo Request from PC1 is flooded on all ports, thus reaching PC2. When PC2 responds, all bridges update their MAC forwarding table using the source MAC address of PC2. Record the amount of time that the ping from PC1 to PC2 is not successful after PC2 has been moved to a different hub.

4. Now connect PC3 to the same hub as PC4.
5. Issue a ping command from PC1 to PC3 that continuously sends ICMP Echo Request packets to PC3:

```
| PC1% ping 10.0.1.31
```

6. Then, generate a single ICMP Echo Request packet from PC3 to PC1 with the command

```
| PC3% ping -c 1 10.0.1.11
```

Now, if you look on PC1, you notice that the ping command is successful again. Explain this outcome and compare it to the outcome of Step 2.

Question 4.B.1)

Include the times that you recorded in Steps 3.

The pings from PC1 to PC2 were successful instantly.

Question 4.B.2)

Explain the outcome of Step 6. That is, explain why the ping issued by PC3 has the effect that the ping commands from PC1 to PC3 (in Step 5) are successful. Compare the outcome with the outcome in Step 3.

The pings became successful without the need of sending the ping from PC3. It took approximately 8 seconds for the ping to start succeeding. We have seen this behaviour multiple times while repeating the exercise.